



ChEESE

Artificial Intelligence and Machine Learning for Geosciences

Physics-informed neural networks (PINNs)

H. Frezat, A. Fournier, L. Seydoux, G. Moguilny

Barcelona, November 7th, 2024



Project funded by EuroHPC under the grant agreement No 101093038.



Trained models and physics

$$f_{\theta}(\mathbf{x}) : \mathbf{x} \rightarrow \mathbf{y}, \quad \operatorname{argmin}_{\theta} \mathcal{L}(f_{\theta}(\mathbf{x}), \mathbf{y})$$

Can I say my model f_{θ} is physically **accurate**?

Trained models and physics

$$f_{\theta}(\mathbf{x}) : \mathbf{x} \rightarrow \mathbf{y}, \quad \operatorname{argmin}_{\theta} \mathcal{L}(f_{\theta}(\mathbf{x}), \mathbf{y})$$

Can I say my model f_{θ} is physically **accurate**?

- ▶ Respect some invariances.
- ▶ Solution of some equation.
- ▶ Stable when used in simulations (for IVPs).

Trained models and physics


$$f_{\theta}(\mathbf{x}) : \mathbf{x} \rightarrow \mathbf{y}, \quad \operatorname{argmin}_{\theta} \mathcal{L}(f_{\theta}(\mathbf{x}), \mathbf{y})$$

Can I say my model f_{θ} is physically **accurate**?

- ▶ Respect some invariances.
- ▶ Solution of some equation.
- ▶ Stable when used in simulations (for IVPs).

How? Any guesses?

Inductive biases



	Technique	Nature	Instances
1	Data manipulation	Both	Augmentation / post processing
2	Loss function	Approx.	PINNs
3	Architecture	Exact	Basic blocks
4	Free parameters	Exact	

Different techniques to physically constrain a model.

Inductive biases

	Technique	Nature	Instances
1	Data manipulation	Both	Augmentation / post processing
2	Loss function	Approx.	PINNs
3	Architecture	Exact	Basic blocks
4	Free parameters	Exact	

Different techniques to physically constrain a model.

Many success stories:

- ▶ Atmospheric and ocean science (Bolton and Zanna, 2019; Beucler, Pritchard, et al., 2021).
- ▶ Weather and climate science (Kashinath et al., 2021; Beucler, Gentine, et al., 2024).
- ▶ Any fluid dynamics-based field.
- ▶ How about solid earth? See Moseley et al. (2020) for the wave equation in complex media

Solving a PDE: the modern way

Let's say we want to solve a non-linear (time-dependent) PDE.

$$u_t + \mathcal{N}(u) = 0$$

$$u(t = 0, x) = g(x)$$

$$u(t, x = \partial\Omega) = c(t)$$

Solving a PDE: the modern way

Let's say we want to solve a non-linear (time-dependent) PDE.

$$u_t + \mathcal{N}(u) = 0$$

$$u(t = 0, x) = g(x)$$

$$u(t, x = \partial\Omega) = c(t)$$

- ▶ Spatial discretization: finite-elements/differences/volumes: mesher, linear algebra solvers, etc.
- ▶ Temporal discretization.
- ▶ Efficient parallelization of the entire process.

Solving a PDE: the modern way

Let's say we want to solve a non-linear (time-dependent) PDE.


$$u_t + \mathcal{N}(u) = 0$$

$$u(t = 0, x) = g(x)$$

$$u(t, x = \partial\Omega) = c(t)$$

- ▶ Spatial discretization: finite-elements/differences/volumes: mesher, linear algebra solvers, etc.
- ▶ Temporal discretization.
- ▶ Efficient parallelization of the entire process.
- ▶ A couple hundred lines of Python?

PINNs: a powerful revolution



Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations



M. Raissi^a, P. Perdikaris^{b,*}, G.E. Karniadakis^a

^a Division of Applied Mathematics, Brown University, Providence, RI, 02912, USA

^b Department of Mechanical Engineering and Applied Mechanics, University of Pennsylvania, Philadelphia, PA, 19104, USA

(Huge success in the scientific community ($> 11k$ citations)).

- ▶ Mesh-free (live in abstract latent space).
- ▶ Data-free (but supervised in the strict sense).
- ▶ Interpolation capabilities / super-resolution.
- ▶ Zero-effort parallelization.

The idea behind PINNs

$$f_{\theta}(t, x)_t + \mathcal{N}[f_{\theta}(t, x)] = 0$$

$$f_{\theta}(t = 0, x) = g(x)$$

$$f_{\theta}(t, x = \partial\Omega) = c(t)$$

A loss for each prior term



- ▶ Boundary conditions loss $\mathcal{L}_{\text{BC}}(f_{\theta}(t, x = \partial\Omega), c(t))$.
- ▶ Initial conditions loss $\mathcal{L}_{\text{IC}}(f_{\theta}(t = 0, x), g(x))$.
- ▶ Equation residual loss $\mathcal{L}_{\text{eq}}(f_{\theta}(t, x)_t + \mathcal{N}[f_{\theta}(t, x)])$.

The total loss is a sum of each:

$$\mathcal{L} = \mathcal{L}_{\text{BC}} + \mathcal{L}_{\text{IC}} + \mathcal{L}_{\text{eq}}.$$

An illustration with Burgers' equation

$$\frac{\partial u}{\partial t} + u \frac{\partial u}{\partial x} - \nu \frac{\partial^2 u}{\partial x^2} = 0$$



Solution of Burgers' equation (Raissi et al., 2019).

- ▶ Let's create and train a neural solver for Burgers' equation in the notebook.

Bonus: hybrid modeling

Sometimes the problem is either **too slow** to solve, or contains **some unknown quantities**:

- ▶ Model complex phenomena not resolved by the equations (microphysics, clouds, etc.).
 - ▶ Accelerate simulations by only resolving large scales (sub-grid modeling).
 - ▶ Probably other instances.
-
- ▶ Keep the physics components with classical numerical methods.
 - ▶ We can explore an example of sub-grid modeling in the bonus notebook.

References



Beucler, Tom, Pierre Gentine, et al. (2024). “Climate-invariant machine learning”. In: *Science Advances* 10.6, eadj7250.



Beucler, Tom, Michael Pritchard, Stephan Rasp, Jordan Ott, Pierre Baldi, and Pierre Gentine (2021). “Enforcing analytic constraints in neural networks emulating physical systems”. In: *Physical Review Letters* 126.9, p. 098302.



Bolton, Thomas and Laure Zanna (2019). “Applications of deep learning to ocean data inference and subgrid parameterization”. In: *Journal of Advances in Modeling Earth Systems* 11.1, pp. 376–399.



Kashinath, Karthik et al. (2021). “Physics-informed machine learning: case studies for weather and climate modelling”. In: *Philosophical Transactions of the Royal Society A* 379.2194, p. 20200093.



Moseley, Ben, Andrew Markham, and Tarje Nissen-Meyer (2020). *Solving the wave equation with physics-informed deep learning*.



Raissi, Maziar, Paris Perdikaris, and George E Karniadakis (2019). “Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations”. In: *Journal of Computational physics* 378, pp. 686–707.